

Leihui Li-FinalHandin-solution

November 4, 2023

1 Basic Data Science in Python - Project

1.1 Leihui Li; leihui@mpe.au.dk

This handin is individual and mandatory to pass the course.

Hand in this .ipynb file, no later than 10/11 kl 23:59.

1.1.1 Project description

For this project, you have almost no limitations. You are going to be working with a [smoke detection dataset](#), to classify a fire alarm. The dataset is located in the data folder, as `$smoke_detection.csv`.

Remember before you begin, to inspect the data - i.e does it need cleaning, what are you working with? Also, remember to separate the dataset in two - the labels ("Fire Alarm"), and the rest of the data.

Concretely, you should use your knowledge of testing, overfitting, and the pros and cons of different models, to predict Fire Alarm (0 for no alarm, 1 for alarm) from the different air quality measures. Try at least two different methods. You must reason about your choice of methods, as well as evaluate your results. Can you visualize your results somehow?

```
[1]: import random, math
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
```

```
[2]: csv_data = pd.read_csv('../data/smoke_detection.csv')
print (csv_data.head)

features, labels = csv_data.iloc[:, 2:-2].to_numpy(), csv_data.iloc[:, -1].
    ↳to_numpy().reshape(-1,1)
print (features.shape, labels.shape)
```

<bound method NDFrame.head of				count	UTC	Temperature[C]
Humidity[%]	TVOC[ppb]	eCO2[ppm]	\			
0	0	1654733331	20.000	57.36	0	400
1	1	1654733332	20.015	56.67	0	400
2	2	1654733333	20.029	55.96	0	400

3	3	1654733334	20.044	55.28	0	400
4	4	1654733335	20.059	54.69	0	400
...
62625	62625	1655130047	18.438	15.79	625	400
62626	62626	1655130048	18.653	15.87	612	400
62627	62627	1655130049	18.867	15.84	627	400
62628	62628	1655130050	19.083	16.04	638	400
62629	62629	1655130051	19.299	16.52	643	400

	Raw H2	Raw Ethanol	Pressure[hPa]	PM1.0	PM2.5	NC0.5	NC1.0	NC2.5	\
0	12306	18520	939.735	0.00	0.00	0.00	0.000	0.000	
1	12345	18651	939.744	0.00	0.00	0.00	0.000	0.000	
2	12374	18764	939.738	0.00	0.00	0.00	0.000	0.000	
3	12390	18849	939.736	0.00	0.00	0.00	0.000	0.000	
4	12403	18921	939.744	0.00	0.00	0.00	0.000	0.000	
...
62625	13723	20569	936.670	0.63	0.65	4.32	0.673	0.015	
62626	13731	20588	936.678	0.61	0.63	4.18	0.652	0.015	
62627	13725	20582	936.687	0.57	0.60	3.95	0.617	0.014	
62628	13712	20566	936.680	0.57	0.59	3.92	0.611	0.014	
62629	13696	20543	936.676	0.57	0.59	3.90	0.607	0.014	

	CNT	Fire Alarm
0	0	0
1	1	0
2	2	0
3	3	0
4	4	0
...
62625	5739	0
62626	5740	0
62627	5741	0
62628	5742	0
62629	5743	0

```
[62630 rows x 16 columns]>
(62630, 12) (62630, 1)
```

1.2 1. Decision Tree

The dataset has multiple features, such as humidity, eCO2, PM*, etc., it is commonly used for decision tree to make the prediction. It is able to model non-linear relationships in the data. This is particularly useful when the underlying data relationships are non-linear, such as the data presented in the smoke dataset. However, when using a decision tree, it can be difficult to capture complex relationships between the data.

1.2.1 1.1 Implementation

```
[3]: # print (features.shape, labels.shape, type(features), type(labels))

from sklearn import tree
np.set_printoptions(precision=3, suppress=True)

X, y = features.copy(), labels.copy()
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4,
    ↪ random_state=42)

print (f"Train data length:{X_train.shape}, Test data length:{X_test.shape}")
x_i = random.randint(0, len(X_test))
print (f"Here is an example: we select \"{x_i}\" th in Test data:")
print (f"feature:{X_train[x_i].tolist()} \nIts label is { y_train[x_i].
    ↪ tolist()}")

# Build decision tree and model
clf_model = tree.DecisionTreeClassifier(
    criterion='entropy',
    splitter = 'random',
    random_state=42,
    max_depth=8,
    #min_impurity_decrease = 0.01,
    min_samples_leaf=1,
    min_samples_split=3)
clf_model = clf_model.fit(X_train, y_train)
# evaluate by score
clf_score_sample, clf_score_test = clf_model.score(X_train,y_train), clf_model.
    ↪ score(X_test,y_test)

print (f"The parameters of this decision tree are as follows:")
print (clf_model.get_params())

print (f"Training score: {clf_score_sample} test score: {clf_score_test}")

print (f"We select the {x_i} th among the test data:")
print (f"ground Truth:\n{X_test[x_i].tolist()} -> {y_test[x_i]}")
print (f"predict Result:\n{clf_model.predict(X_test[x_i].reshape(1, -1))}")
```

Train data length:(37578, 12), Test data length:(25052, 12)

Here is an example: we select "12017" th in Test data:

feature:[21.971, 49.32, 39.0, 455.0, 12541.0, 19629.0, 939.817, 0.0, 0.0, 0.0, 0.001, 0.001]

Its label is [0]

The parameters of this decision tree are as follows:

{'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'entropy', 'max_depth': 8, 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease': 0.0,

```
'min_samples_leaf': 1, 'min_samples_split': 3, 'min_weight_fraction_leaf': 0.0,
'random_state': 42, 'splitter': 'random'}
Training score: 0.9736015753898558 test score: 0.9741737186651764
We select the 12017 th among the test data:
ground Truth:
[29.39, 40.4, 53.0, 400.0, 12842.0, 20733.0, 937.571, 1.98, 2.06, 13.66, 2.13,
0.048] -> [0]
predict Result:
[0]
```

According to the result of the decision tree, the accuracy for the test dataset is about 97%, and the trained model also has 97% accuracy. The overfitting does not seem to have happened here.

1.3 1.2 Visualization

```
[4]: text_representation = tree.export_text(clf_model)
print(text_representation)

import sklearn
plt.figure(figsize=(25,20)) # Resize figure
sklearn.tree.plot_tree(clf_model, filled=True)
plt.show()


|--- feature_5 <= 19886.50
|   |--- feature_2 <= 27564.26
|   |   |--- feature_0 <= 47.93
|   |   |   |--- feature_0 <= 31.99
|   |   |   |   |--- feature_1 <= 40.65
|   |   |   |   |   |--- feature_6 <= 934.13
|   |   |   |   |   |   |--- feature_9 <= 1422.27
|   |   |   |   |   |   |   |--- feature_6 <= 931.09
|   |   |   |   |   |   |   |   |--- class: 1
|   |   |   |   |   |   |   |   |--- feature_6 > 931.09
|   |   |   |   |   |   |   |   |--- class: 0
|   |   |   |   |   |   |   |   |--- feature_9 > 1422.27
|   |   |   |   |   |   |   |   |--- class: 1
|   |   |   |   |   |   |   |   |--- feature_6 > 934.13
|   |   |   |   |   |   |   |   |--- class: 0
|   |   |   |   |   |--- feature_1 > 40.65
|   |   |   |   |   |--- feature_6 <= 937.36
|   |   |   |   |   |--- feature_2 <= 1113.46
|   |   |   |   |   |--- feature_8 <= 1893.77
|   |   |   |   |   |   |--- class: 0
|   |   |   |   |   |   |--- feature_8 > 1893.77
|   |   |   |   |   |   |   |--- class: 1
|   |   |   |   |   |   |   |--- feature_2 > 1113.46
|   |   |   |   |   |   |   |   |--- class: 0
|   |   |   |   |   |--- feature_6 > 937.36
|   |   |   |   |   |--- feature_9 <= 4.31
```

```

| | | | | | | | |--- feature_2 <= 235.64
| | | | | | | | |--- class: 0
| | | | | | | | |--- feature_2 > 235.64
| | | | | | | | |--- class: 1
| | | | | | | | |--- feature_9 > 4.31
| | | | | | | | |--- class: 1
| | | |--- feature_0 > 31.99
| | | |--- feature_6 <= 932.52
| | | |--- class: 1
| | | |--- feature_6 > 932.52
| | | |--- class: 0
| | |--- feature_0 > 47.93
| | |--- class: 0
| |--- feature_2 > 27564.26
| |--- class: 0
|--- feature_5 > 19886.50
| |--- feature_5 <= 20696.71
| | |--- feature_5 <= 20284.74
| | | |--- feature_0 <= 17.04
| | | |--- feature_0 <= 10.54
| | | | |--- feature_2 <= 1774.15
| | | | |--- feature_0 <= 0.12
| | | | |--- feature_1 <= 25.74
| | | | |--- class: 0
| | | | |--- feature_1 > 25.74
| | | | |--- class: 1
| | | | |--- feature_0 > 0.12
| | | | |--- feature_1 <= 52.22
| | | | |--- class: 1
| | | | |--- feature_1 > 52.22
| | | | |--- class: 1
| | | | |--- feature_2 > 1774.15
| | | | |--- class: 0
| | | | |--- feature_0 > 10.54
| | | | |--- class: 0
| | | |--- feature_0 > 17.04
| | | |--- feature_2 <= 2265.57
| | | |--- feature_5 <= 20213.41
| | | |--- feature_2 <= 25.99
| | | |--- feature_5 <= 20189.29
| | | |--- class: 0
| | | |--- feature_5 > 20189.29
| | | |--- class: 1
| | | |--- feature_2 > 25.99
| | | |--- feature_0 <= 24.11
| | | |--- class: 1
| | | |--- feature_0 > 24.11
| | | |--- class: 0

```

```
| | | | | |--- feature_5 > 20213.41
| | | | | | |--- feature_2 <= 565.76
| | | | | | |--- feature_6 <= 938.53
| | | | | | |--- class: 1
| | | | | | |--- feature_6 > 938.53
| | | | | | |--- class: 1
| | | | | | |--- feature_2 > 565.76
| | | | | | |--- class: 0
| | | | | |--- feature_2 > 2265.57
| | | | | |--- class: 0
| | | | | |--- feature_5 > 20284.74
| | | | | |--- feature_6 <= 935.92
| | | | | |--- class: 1
| | | | | |--- feature_6 > 935.92
| | | | | |--- feature_4 <= 13265.38
| | | | | |--- class: 0
| | | | | |--- feature_4 > 13265.38
| | | | | |--- feature_6 <= 937.37
| | | | | |--- class: 0
| | | | | |--- feature_6 > 937.37
| | | | | |--- class: 1
| | | | | |--- feature_5 > 20696.71
| | | | | |--- feature_6 <= 935.73
| | | | | |--- class: 1
| | | | | |--- feature_6 > 935.73
| | | | | |--- feature_3 <= 427.35
| | | | | |--- class: 0
| | | | | |--- feature_3 > 427.35
| | | | | |--- feature_10 <= 2.37
| | | | | |--- class: 0
| | | | | |--- feature_10 > 2.37
| | | | | |--- feature_9 <= 15.59
| | | | | |--- feature_3 <= 452.08
| | | | | |--- class: 0
| | | | | |--- feature_3 > 452.08
| | | | | |--- class: 1
| | | | | |--- feature_9 > 15.59
| | | | | |--- class: 0
```



1.4 2. Neural Networks

1.5 2.1 Build the regression model using PyTorch

```
[37]: import torch, time
from tqdm.notebook import tqdm
from torch.utils.data import Dataset, DataLoader
import torch.nn as nn
import torch.nn.functional as F
from IPython import display
from sklearn.metrics import accuracy_score

def accuracy_fn(y_true, y_pred):
    correct = torch.eq(y_true, y_pred).sum().item() # torch.eq() calculates_
    ↳where two tensors are equal
    acc = (correct / len(y_pred)) * 100
    return acc

train_loss_list, test_loss_list, ac_list = [], [], []
```

```

def run_epoch(model, optimizer, data_loader, loss_func, device, results,
    ↪score_funcs, prefix="", desc=None):
    """
    model -- the PyTorch model / "Module" to run for one epoch
    optimizer -- the object that will update the weights of the network
    data_loader -- DataLoader object that returns tuples of (input, label) ↪
    ↪pairs.
    loss_func -- the loss function that takes in two arguments, the model ↪
    ↪outputs and the labels, and returns a score
    device -- the compute lodation to perform training
    score_funcs -- a dictionary of scoring functions to use to evaluate the ↪
    ↪performance of the model
    prefix -- a string to pre-fix to any scores placed into the _results_ ↪
    ↪dictionary.
    desc -- a description to use for the progress bar.
    """
    running_loss = []
    y_true = []
    y_pred = []
    start = time.time()
    for inputs, labels in tqdm(data_loader, desc=desc, leave=False):
        #Move the batch to the device we are using.
        inputs = moveTo(inputs, device)
        labels = moveTo(labels, device)

        y_hat = model(inputs) #this just computed  $f_{\theta}(x(i))$ 

        # Compute loss.
        loss = loss_func(y_hat, labels)

        if model.training:
            loss.backward()
            optimizer.step()
            optimizer.zero_grad()

        #Now we are just grabbing some information we would like to have
        running_loss.append(loss.item())
        if score_funcs is not None:
            if len(score_funcs) > 0 and isinstance(labels, torch.Tensor):
                #moving labels & predictions back to CPU for computing ↪
    ↪storing predictions
                labels = labels.detach().cpu().numpy()
                y_hat = y_hat.detach().cpu().numpy()
                #add to predictions so far
                y_true.extend(labels.tolist())
                y_pred.extend(y_hat.tolist())

```



```

#end training epoch
end = time.time()

y_pred = np.asarray(y_pred)
if len(y_pred.shape) == 2 and y_pred.shape[1] > 1: #We have a
↳classification problem, convert to labels
    y_pred = np.argmax(y_pred, axis=1)
    #Else, we assume we are working on a regression problem

acc = accuracy_score(y_true, y_pred)
#    print (f'acc:{acc}')

#    print(f'loss:{np.mean(running_loss)}')
if prefix == 'train':
    train_loss_list.append(np.mean(running_loss))
elif prefix == 'test':
    ac_list.append(acc)
    test_loss_list.append(np.mean(running_loss))
display.clear_output(wait=True)
plt.plot(train_loss_list, label = 'train loss')
plt.plot(test_loss_list, label = 'loss loss')
plt.plot(ac_list, label = 'accuracy')
plt.legend()
plt.show()

results[prefix + " loss"].append( np.mean(running_loss) )
if score_funcs is not None:
    for name, score_func in score_funcs.items():
        try:
            results[prefix + " " + name].append( score_func(y_true, y_pred)
↳)
        except:
            results[prefix + " " + name].append(float("NaN"))
return end-start #time spent on epoch

def train_simple_network(model, loss_func, train_loader, test_loader=None,
↳score_funcs=None,
                            epochs=20, device="cpu", checkpoint_file=None, _lr = 0.
↳01):
    to_track = ["epoch", "total time", "train loss"]
    if test_loader is not None:
        to_track.append("test loss")
    if score_funcs is not None:
        for eval_score in score_funcs:
            to_track.append("train " + eval_score )
            if test_loader is not None:
                to_track.append("test " + eval_score )

```

```

total_train_time = 0 #How long have we spent in the training loop?
results = {}
#Initialize every item with an empty list
for item in to_track:
    results[item] = []

#SGD is Stochastic Gradient Decent.
optimizer = torch.optim.SGD(model.parameters(), lr=_lr)
# optimizer = torch.optim.Adam(model.parameters(), lr=_lr)
#Place the model on the correct compute resource (CPU or GPU)
model.to(device)
for epoch in tqdm(range(epochs), desc="Epoch"):
    model = model.train() #Put our model in training mode

    total_train_time += run_epoch(model, optimizer, train_loader,
↳loss_func, device, results, score_funcs, prefix="train", desc="Training")

    results["total time"].append( total_train_time )
    results["epoch"].append( epoch )

    if test_loader is not None:
        model = model.eval()
        with torch.no_grad():
            run_epoch(model, optimizer, test_loader, loss_func, device,
↳results, score_funcs, prefix="test", desc="Testing")

    if checkpoint_file is not None:
        torch.save({
            'epoch': epoch,
            'model_state_dict': model.state_dict(),
            'optimizer_state_dict': optimizer.state_dict(),
            'results' : results
        }, checkpoint_file)

    return pd.DataFrame.from_dict(results)

class RegressionDataset(Dataset):
    def __init__(self, X, y):
        # super(Simple1DRegressionDataset, self).__init__()
        self.X = X.reshape(-1,12)
        self.y = y.reshape(-1,1)

    def __getitem__(self, index):
        return torch.tensor(self.X[index,:], dtype=torch.float32), torch.
↳tensor(self.y[index], dtype=torch.float32)

```

```

def __len__(self):
    return self.X.shape[0]

def moveTo(obj, device):
    """
    obj: the python object to move to a device, or to move its contents to a
    ↪ device
    device: the compute device to move objects to
    """
    if isinstance(obj, list):
        return [moveTo(x, device) for x in obj]
    elif isinstance(obj, tuple):
        return tuple(moveTo(list(obj), device))
    elif isinstance(obj, set):
        return set(moveTo(list(obj), device))
    elif isinstance(obj, dict):
        to_ret = dict()
        for key, value in obj.items():
            to_ret[moveTo(key, device)] = moveTo(value, device)
        return to_ret
    elif hasattr(obj, "to"):
        return obj.to(device)
    else:
        return obj

```

```

[38]: ## Dataset preparation
device = torch.device("cuda")
# For this dataset, the index number and the UTC may not be helpful, so 12
↪ features could be enough
in_features = 12
out_features = 2
node_num = 10
model = nn.Sequential(
    nn.Linear(in_features, node_num), #hidden layer
    nn.Tanh(), #activation
    nn.Linear(node_num, node_num), #hidden layer
    nn.Tanh(), #activation
    nn.Linear(node_num, node_num), #hidden layer
    nn.Tanh(), #activation
    nn.Linear(node_num, node_num), #hidden layer
    nn.Tanh(), #activation
    nn.Linear(node_num, out_features), #output layer
)
torch.manual_seed(3407)
X, y = features.copy(), labels.copy()
print (X.shape, y.shape)

```

```

dataset = torch.utils.data.TensorDataset(
    torch.tensor(X, dtype=torch.float32).to(device),
    torch.tensor(y.flatten(), dtype=torch.long).to(device))
train_ratio = 0.8
num_of_train_rows = int(X.shape[0]*train_ratio)
num_of_test_rows = int(math.ceil(X.shape[0]*(1-train_ratio)))
print (f"{num_of_train_rows} and {num_of_test_rows} == {X.shape[0]}")

generator1 = torch.Generator().manual_seed(3407)
train_data, test_data = torch.utils.data.random_split(
    dataset,
    lengths=[num_of_train_rows, num_of_test_rows],
    generator=generator1)

```


(62630, 12) (62630, 1)
50104 and 12526 == 62630

```

[39]: my_epoch = 100
my_lr = 0.01
my_batch_size = 64

training_loader = DataLoader(train_data, shuffle=True, batch_size=my_batch_size)
testing_loader = DataLoader(test_data, shuffle=True, batch_size=512)
## Training
loss_func = nn.CrossEntropyLoss()
checkpoint_name = f'best_epoch_{my_epoch}_lr_{my_lr}_bt_{my_batch_size}'
results = train_simple_network(model,
                                loss_func, training_loader, testing_loader,
                                ↪score_funcs={'Accuracy': accuracy_score},
                                epochs=my_epoch, device=device,
                                ↪checkpoint_file=f'{checkpoint_name}.pt', _lr=my_lr)


```



```
[40]: plt.plot(results['train loss'], label='train loss')
plt.plot(results['test loss'], label='test loss')
plt.legend()
plt.savefig(f'{checkpoint_name}.png')
print (results)
```

	epoch	total time	train loss	test loss	train Accuracy	test Accuracy
0	0	1.634307	0.597235	0.586566	0.705513	0.719863
1	1	3.322193	0.583265	0.588863	0.724333	0.720661
2	2	4.964144	0.584900	0.585342	0.729323	0.729443
3	3	6.666305	0.583200	0.582978	0.733175	0.731838
4	4	8.319234	0.580986	0.582121	0.732736	0.730720
..
95	95	179.357881	0.574404	0.575785	0.734353	0.732476
96	96	181.339846	0.574419	0.577088	0.734353	0.732476
97	97	183.198995	0.574379	0.576437	0.734353	0.732476
98	98	184.969255	0.574340	0.575432	0.734353	0.732476
99	99	186.883230	0.574353	0.576171	0.734353	0.732476

[100 rows x 6 columns]



```
[41]: checkpoint = torch.load(f'{checkpoint_name}.pt')
model.load_state_dict(checkpoint['model_state_dict'])
model.to(device)
model.eval()

def pref(_x):
    with torch.no_grad():
        _x = torch.as_tensor(_x, dtype=torch.float32).to(device)
        Y_logits = model(_x).clone().detach()
        print (Y_logits, type(Y_logits))
        y_hat = np.argmax(Y_logits.cpu())
        print (y_hat)

user_x_0 = torch.tensor(
    np.array([27.089, 55.94, 19, 400, 13010, 19919, 939.742, 0.15, 0.15, 1.02, 0.158, 0.
    ↪004]),
    dtype =torch.float32).to(device) # 0
user_x_1 = torch.tensor(
    np.array([-1.144, 48.22, 1325, 408, 12986, 19394, 938.809, 2.08, 2.16, 14.33, 2.235, 0.
    ↪05]),
    dtype =torch.float32).to(device) # 1
```

```
pref(user_x_0)
pref(user_x_1)
```


```
tensor([-0.3029,  0.7471], device='cuda:0') <class 'torch.Tensor'>
tensor(1)
tensor([-0.3029,  0.7471], device='cuda:0') <class 'torch.Tensor'>
tensor(1)
```


Your thoughts and reflections here In this section the neural network for smoke detection is designed and implemented. It is a classification problem, so the loss function is chosen as **CrossEntropyLoss**. I have tried several training methods where the structure of the network is consistent and the data set is split again and the number of epochs is changed.

For the training with the same epoch, the result including training loss and test loss is shown below. The reason of showing the difference between two training results, is probably the **stratified sampling for class imbalance**. More specifically, the fire alarm (label == 1) account for 44757 (44757/62631 = 71.4%) in our dataset, therefore the model learns to simply reiterate the most common class label even if we randomly split the dataset. The trained model can always predict fire alarm (1) with high accuracy for the test data. This also means that the designed network is overfitting for our data set. The accuracy is always ~75%, and it is hard to increase, even though I have tried adding more layers, changing the loss function, and increasing/decreasing the `batch_size`. I realise there is a training bottleneck. I would really like to reproduce the 99% accuracy presented in the link.

In addition, according to the introduction of the data generation method, the input vector is time-dependent, which means that the data are sequences. Therefore, the RNN network might be the better solution.

```
[42]: from PIL import Image
      from IPython.display import display
      img_0_epoch_100 = Image.open('0-epoch-100.png')
      img_1_epoch_100 = Image.open('1-epoch-100.png')
      display(img_0_epoch_100),display(img_1_epoch_100)
```







[42]: (None, None)

Second, I tried the training with more epochs, 300 and 500 epochs, while the dataset is the same, which means that the split did not run again. The results of this training are shown below. According to the results, I would like to say that the **100 of epoch** is enough to use for the current structure of the network.

```
[43]: img_epoch_300 = Image.open('1-epoch-300.png')
      img_epoch_500 = Image.open('0-epoch-500.png')
      display(img_epoch_300),display(img_epoch_500)
```





[43]: (None, None)

I tried more tricks on training, the point is that I have to adjust more parameters such as the network layers, epochs and learning rate. So far, for this network, I found that the more epochs the network trains, the worse overfitting will be.